APAI Lab03: DNN Shrinking and Quantization

Davide Nadalini, Lorenzo Lamberti, Alberto Dequino, Luca Bompani, Francesco Conti.

(University of Bologna)

d.nadalini@unibo.it lorenzo.lamberti@unibo.it
alberto.dequino@unibo.it

In this Hands-on session:

A first-time user of Pytorch framework will learn how to :

- shrink a NN, by acting on the number of layers, channels, or stride factor
- Quantize a NN down to 2 bits
- Use Netron to visualize a ONNX representation of a CNN

Tasks:

1. Load model's trained weights of LAB1;
2. Reduce network's size under 5 MMAC;
3. Re-train the reduced network and verify network's accuracy;
4. Quantize with QuantLab;
5. Export Onnx and analyze the float32 and quantized models with Netron.

All the details about the tasks are explained in the pdf document attached.

How to deliver the assignment

- Use Virtuale platform to load your file
- update only the .ipynb file, **named as follows**:

  LAB1_APAI_yourname.ipynb

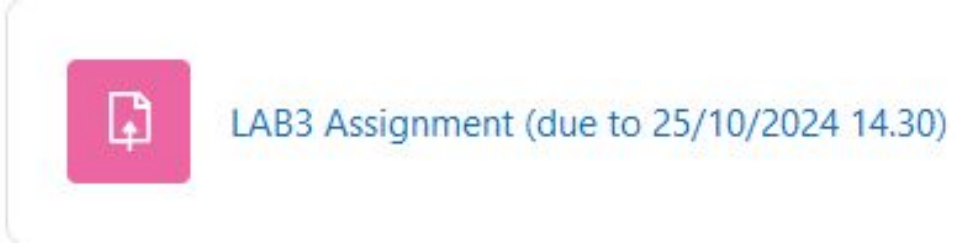

Important: the notebook must be pre-run by you. Outputs must be correct and visible when you download it.

LAB1 DEADLINE:

25/10/2024 at 14:30

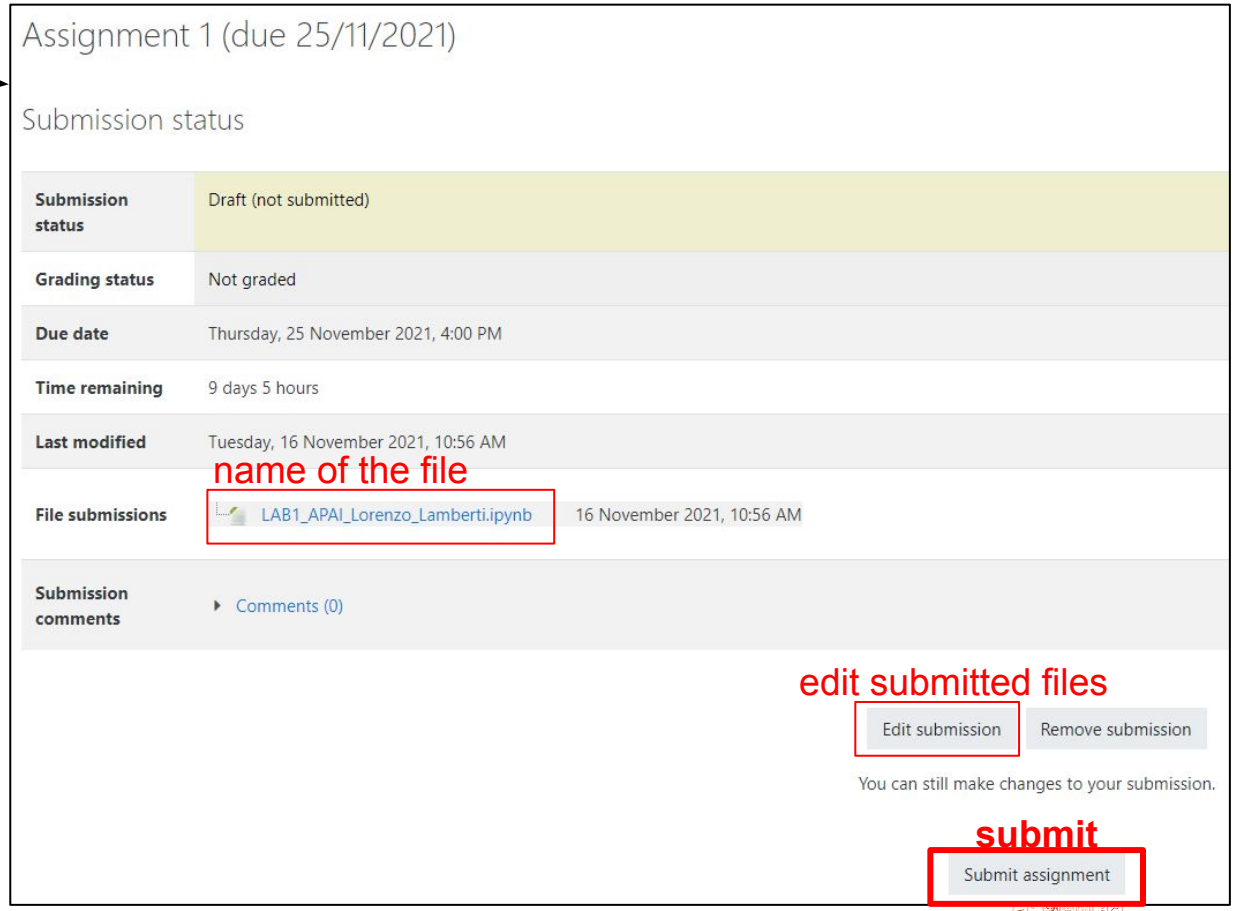

How to download the .ipynb file

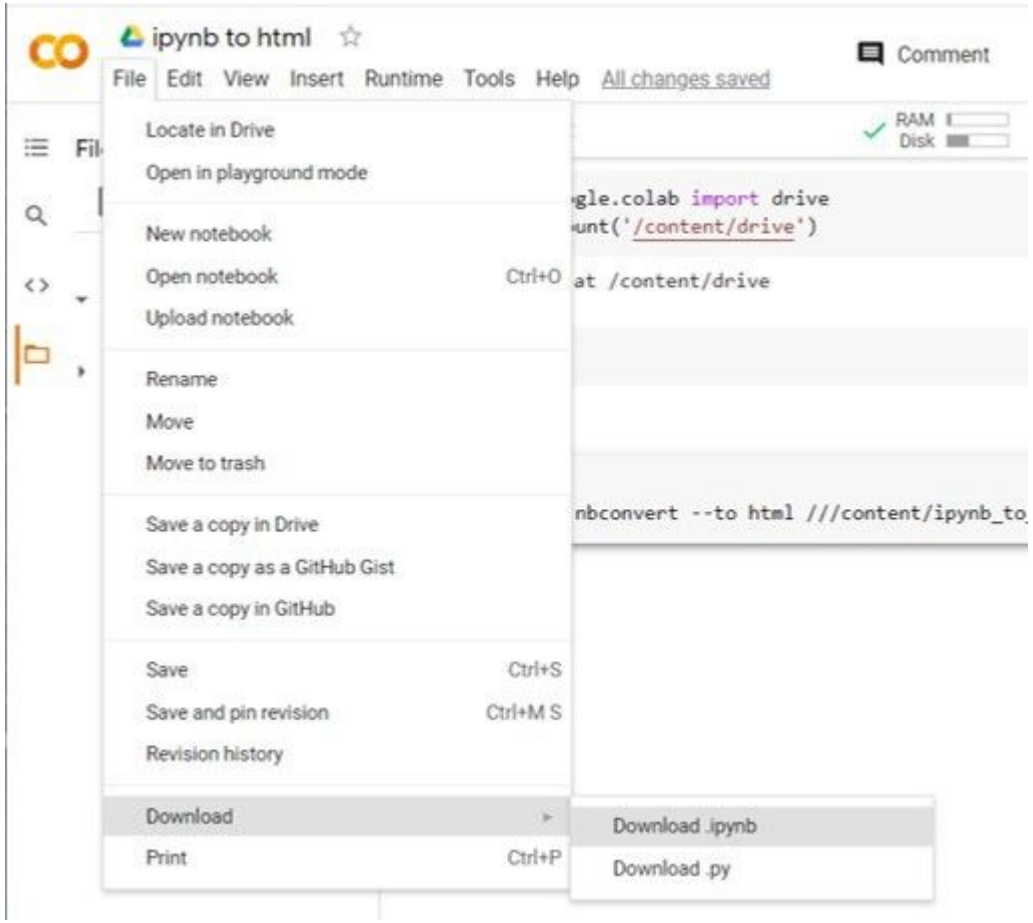

Setup

1

Links to COLAB exercise:

GitHub

Solution is coming after the deadline

→ **Open github link**

2

APAI24-LAB03-DNN-Shrinking-and-Quantization

Guidelines:

1. Start by reading the slides;
2. Then read the assignment;
3. Now complete the assignment: colab.

→ **Open Colab Jupyter notebook**

3

Open in Colab

LAB03 APAI: DNN shrinking & quantization

Credits: *Davide Nadalini, Lorenzo Lamberti, Luca Bompani, Alberto Dequino, Francesco Conti. (University of Bologna)*

Contacts: lorenzo.lamberti@unibo.it, d.nadalini@unibo.it, alberto.dequino@unibo.it, luca.bompani5@unibo.it

→ **Open In COLAB to modify it !**

4

→ **create your own copy in COLAB to modify it !**

more details on step [4] of the setup

IMPORTANT:

Create your own copy of the COLAB notebook!

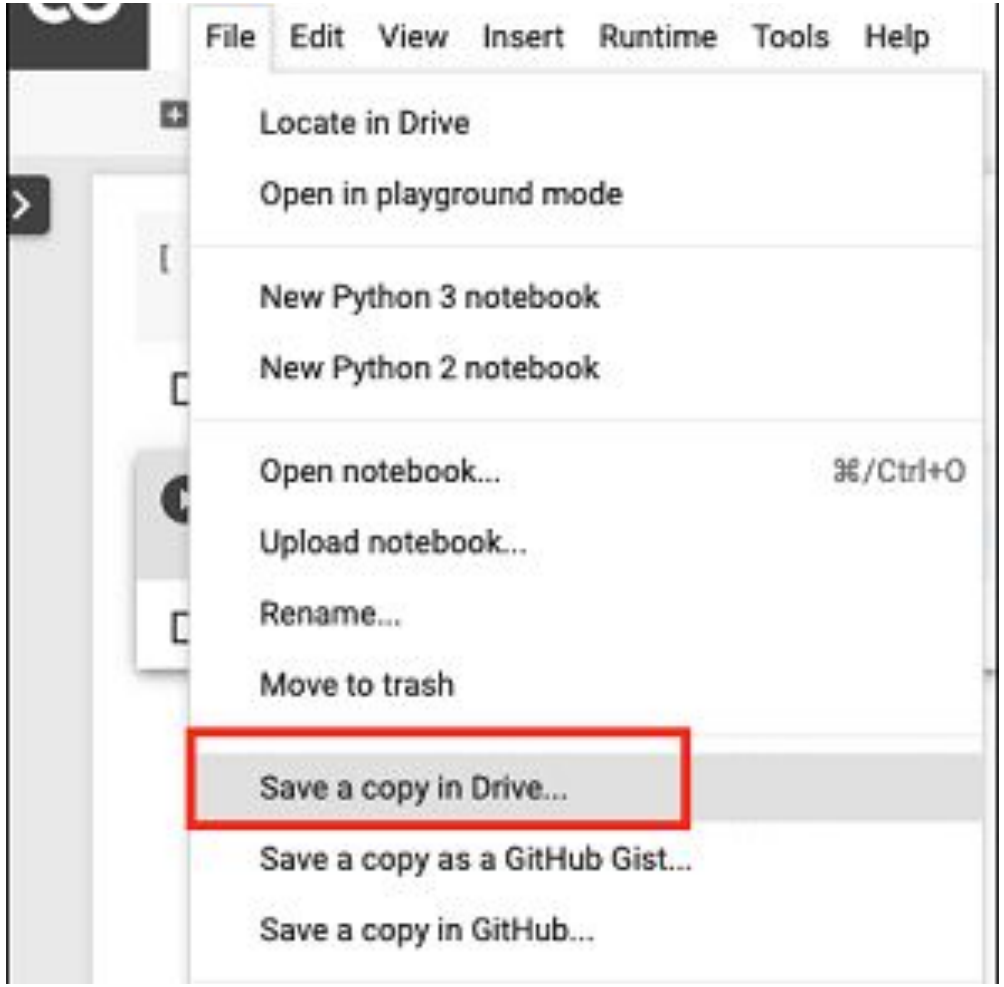

Others:

- Activate/deactivate GPU: Runtime -> Change runtime type
- **Note:** If you use for too much time the GPU, your account will be limited to CPU for 24h.

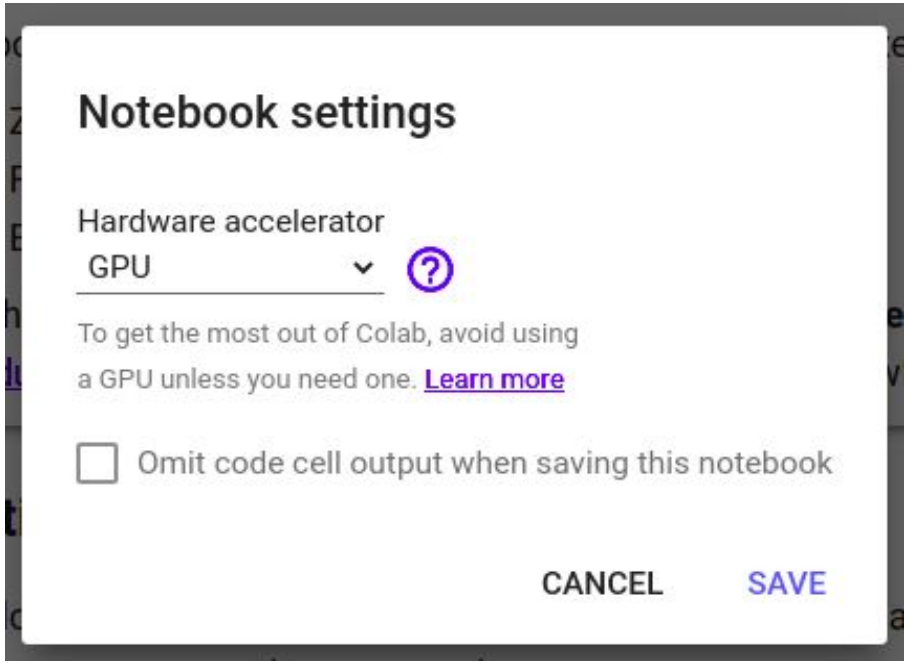

ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

A bit of theory

Data Types and Quantization

IEEE 754 32-bit Floating Point format

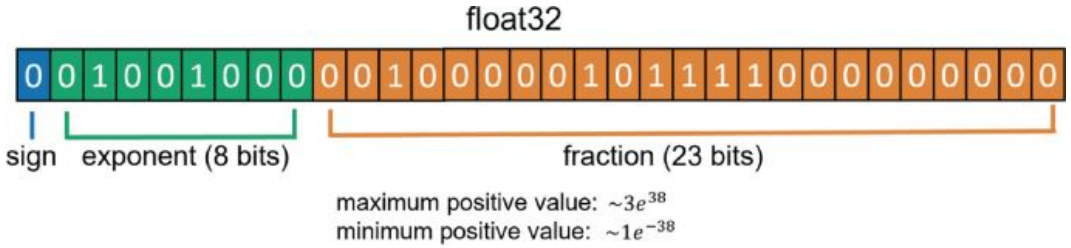

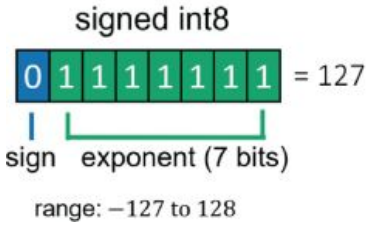

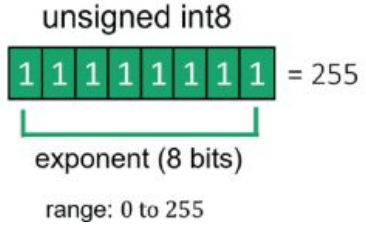

Training time: FP32 for numerical accuracy! However, 1 element of activation / weights = 4 bytes (memory hungry for inference!)

Deployment: quantization! Smaller number of bits for a single element (e.g., 8 bits = 1 byte), but less numerical precision!!

Remember Lab1? **Data stored in 8 bits** (instead of 16 or 32) easily **overflow** or cannot represent **values outside range**! One possible solution: **quantization-aware training**!

Sources: https://en.wikipedia.org/wiki/Single-precision_floating-point_format
https://link.springer.com/chapter/10.1007/978-3-031-24538-1_1

ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA

The LAB starts now !